

Muhammad Hanzla

Full-Stack Developer | AI-Integrated Applications

+92 326 6099884 | hanzlasabir658@gmail.com | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Full-stack developer specializing in AI-integrated applications, working across React, Next.js, Node.js, and FastAPI. Experienced in building RAG pipelines, multi-agent systems with LangGraph, and workflow automation with n8n. B.S. Software Engineering, FAST NUCES. Comfortable owning a feature from database design through to a polished, production-ready interface.

EDUCATION

FAST NUCES, Faisalabad -- B.S. Software Engineering

Aug 2022 - Jun 2026

TECHNICAL SKILLS

Languages: JavaScript, TypeScript, Python, SQL

Frontend: React, Next.js, Redux, Tailwind CSS

Backend: Node.js, Express.js, FastAPI, RESTful APIs

Database: MongoDB, MySQL, PostgreSQL, Prisma ORM, Redis

AI & LLM Integration: LangChain, LangGraph, RAG Pipelines, Prompt Engineering, Multi-Agent Systems

Agentic AI & Automation: Agent Orchestration, Tool Use, n8n Workflow Automation

DevOps & Cloud: Docker, Kubernetes, AWS, CI/CD, GitHub Actions

AI-Assisted Development: Claude, Antigravity

PROFESSIONAL EXPERIENCE

Generative AI Engineer Intern -- Elevvo Pathways

Sep 2025 - Dec 2025

Cairo, Egypt (Remote)

- Designed and deployed Mira, a production clinical AI chatbot, integrating the Cohere API into a multi-stage RAG pipeline with custom prompt engineering -- boosting contextual response relevance by 40% over baseline.
- Elevated retrieval precision by 35% via parent-child chunking and Cohere Rerank v3; optimised inference throughput with hybrid decoding and response caching on Streamlit Cloud for real-time, low-latency responses.

Full Stack Developer Intern -- Nexium

Aug 2025 - Sep 2025

Lahore, Pakistan (Remote)

- Shipped a production e-commerce platform with Redux and Role-Based Access Control (RBAC), cutting auth-related issues by 30%; architected CI/CD pipelines with Appium-automated mobile testing, accelerating releases by 50%.
- Provisioned AWS environments with Terraform, reducing environment setup from days to minutes.

PROJECTS

PROJECTIFY -- Full-Stack FYP Management Platform

Aug 2025 - May 2026

[GitHub](#) | [Live Demo](#)

- Architected a multi-campus FYP lifecycle platform (proposal submission, AI plagiarism detection, panel evaluation, milestone grading) using Next.js 14, Prisma ORM, PostgreSQL, and Socket.IO + Redis for real-time collaboration.
- Built an AI similarity pipeline (Cohere + Pinecone) and an AI panel suggestion engine; engineered a zero-egress Cloudflare R2 client and in-process Gmail SMTP meeting daemon.

HEALIX -- AI-Powered CI/CD Auto-Healing System

Feb 2026 - Mar 2026

[GitHub](#) | [Live Demo](#)

- Cut MTTR by 60% by building an autonomous agent that detects failed GitHub Actions via webhooks, diagnoses root causes with Gemini AI, generates patches via Groq, and opens fix PRs automatically.

SCIO AI -- Autonomous Multi-Agent Research Platform

Jan 2026 - Feb 2026

[GitHub](#) | [Live Demo](#)

- Reduced manual research time by 70% via a LangGraph multi-agent pipeline (Researcher, Writer, Critic) producing citation-backed reports; integrated Tavily search and Groq/OpenRouter LLM routing on Next.js 14 and FastAPI.
- Built a responsive Next.js interface with live agent progress tracking and streaming report rendering.

AXIOS -- Multi-Mode AI Learning Workspace

Mar 2026 - Apr 2026

[GitHub](#) | [Live Demo](#)

- Delivered a five-mode AI platform (Chat, Explain, Quiz, Summarise, Flashcards) with Hybrid Search (FAISS + BM25), Cross-Encoder reranking, real-time SSE streaming, and multi-LLM fallback routing.
- Designed a responsive, component-driven Tailwind frontend across all five modes, including live streaming chat.